

Super Convergence: Very Fast Training of Residual Networks Using Large Learning Rates

ICLR Reproducibility Challenge

Keivaun Waugh
University of Texas at Austin
keivaunwaugh@gmail.com

Paul Choi
University of Texas at Austin
choipaul96@gmail.com

Abstract

In this paper, we aim to evaluate the reproducibility of the experiments detailed in the paper “Super Convergence: Very Fast Training of Residual Networks Using Large Learning Rates” [?]

1. Introduction

Paul

1.1. Target Questions

2. Method

Due to time and hardware constraints (lack of multiple GPUs), we were unable to replicate all of the experiments that were listed in the paper. Therefore, we had to choose which experiments from which we could glean the most.

2.1. Implementation Details

Discuss things like batch size that we were limited by because we only had one GPU to run things on at a time (8GB of memory)

3. Experiments

Keivaun

3.1. Methodology

4. Conclusion

Paul

4.1. Cost of Reproduction

What cost in terms of resources (computation, time, people, development effort, communication with the authors).